


Lecture 20: Topo-Sort and Dijkstra's Greedy Idea

◆ Items on Today's Lunch Menu:

- ⇒ Topological Sort (ver. 1 & 2): Gunning for linear time...
- ⇒ Finding Shortest Paths
 - ◆ Breadth-First Search
 - ◆ Dijkstra's Method: Greed is good!

◆ Covered in Chapter 9 in the textbook

Graph Algorithm #1: Topological Sort




Problem: Find an order in which all these courses can be taken.

Example: 142 143 378
 370 321 341 322
 326 421 401

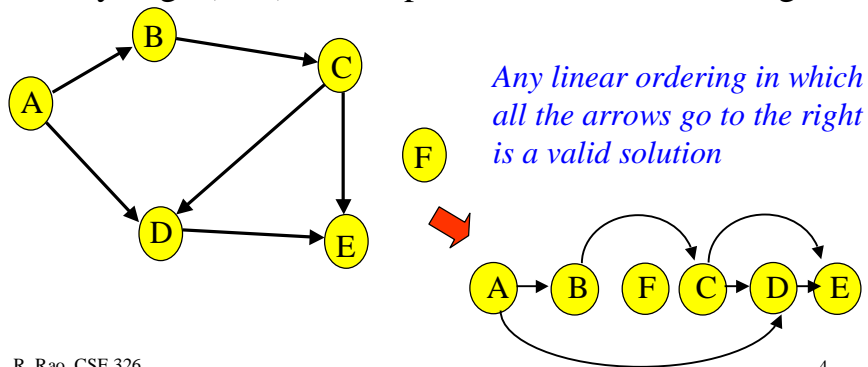
Topological Sort Definition

Topological sorting problem: given digraph $G = (V, E)$, find a linear ordering of vertices such that:
for all edges (v, w) in E , v precedes w in the ordering



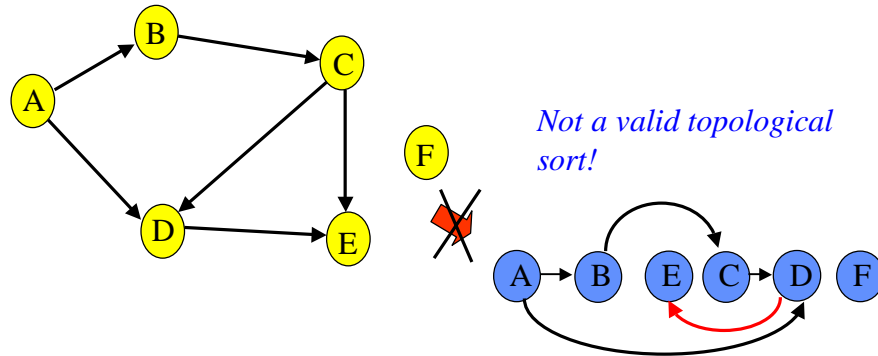
Topological Sort

Topological sorting problem: given digraph $G = (V, E)$, find a linear ordering of vertices such that:
for any edge (v, w) in E , v precedes w in the ordering



Topological Sort


Topological sorting problem: given digraph $G = (V, E)$, find a linear ordering of vertices such that:
for any edge (v, w) in E , v precedes w in the ordering



Topological Sort Algorithm

Step 1: Identify vertices that have no incoming edge


- The “**in-degree**” of these vertices is zero



Topological Sort Algorithm

Step 1: Identify vertices that have no incoming edge

- If *no such edges*, graph has cycles (cyclic graph)




Example of a cyclic graph:
No vertex of in-degree 0

Topological Sort Algorithm

Step 1: Identify vertices that have no incoming edges


- Select one such vertex

Select




Topological Sort Algorithm

Step 2: Delete this vertex of in-degree 0 and all its outgoing edges from the graph. Place it in the output.




Topological Sort Algorithm

Repeat Steps 1 and Step 2 until graph is empty




Topological Sort Algorithm

Repeat Steps 1 and Step 2 until graph is empty



Topological Sort Algorithm

Repeat Steps 1 and Step 2 until graph is empty



Topological Sort Algorithm


Repeat Steps 1 and Step 2 until graph is empty

Final Result:



Summary of Topo-Sort Algorithm #1

1. Store each vertex's **In-Degree** (# of incoming edges) in an array
2. While there are vertices remaining:
 - ⇒ Find a vertex with In-Degree zero and output it
 - ⇒ Reduce In-Degree of all vertices adjacent to it by 1
 - ⇒ Mark this vertex (In-Degree = -1)



0	A	→	B	→	D	/
1	B	→	C	/	/	/
1	C	→	D	→	E	/
2	D	→	E	/	/	/
2	E	/	/	/	/	/
0	F	/	/	/	/	/

In-Degree array Adjacency list

Topological Sort Algorithm #1: Analysis

For input graph $G = (V, E)$, Run Time = ?

Break down into total time required to:

§ Initialize In-Degree array:

$O(|E|)$

§ Find vertex with in-degree 0:

$|V|$ vertices, each takes $O(|V|)$ to search In-Degree array.

Total time = $O(|V|^2)$

§ Reduce In-Degree of all vertices adjacent to a vertex:

$O(|E|)$

§ Output and mark vertex:

$O(|V|)$

Total time = $O(|V|^2 + |E|)$ **Quadratic time!**

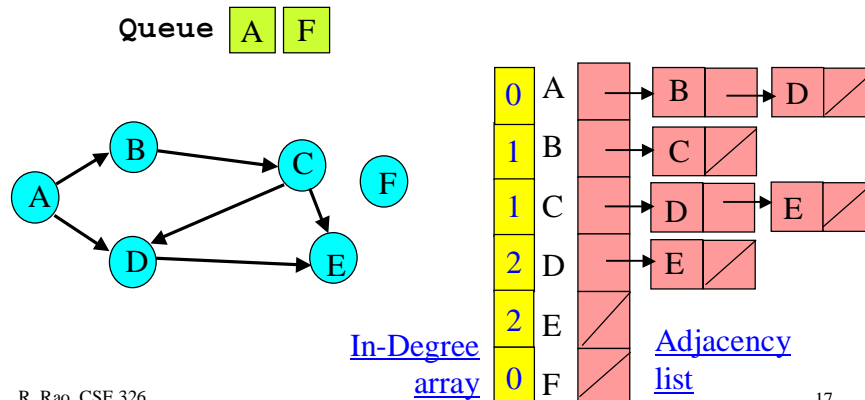
Can we do better than quadratic time?

Problem:

Need a faster way to find vertices with in-degree 0
instead of searching through entire in-degree array

Topological Sort (Take 2)

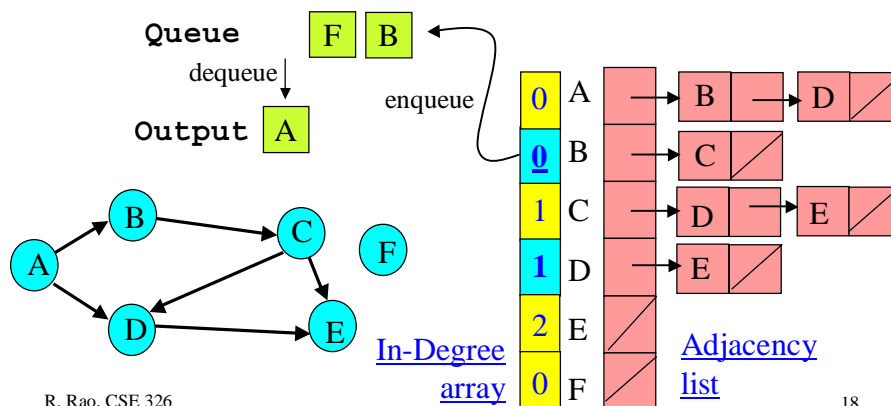
Key idea: Initialize and maintain a *queue (or stack)* of vertices with In-Degree 0



17

Topological Sort (Take 2)


After each vertex is output, when updating In-Degree array, *enqueue any vertex whose In-Degree has become zero*



18

Topological Sort Algorithm #2

1. Store each vertex's **In-Degree** in an array
2. Initialize a queue with all in-degree zero vertices
3. While there are vertices remaining in the queue:
 - ⇒ Dequeue and output a vertex
 - ⇒ Reduce In-Degree of all vertices adjacent to it by 1
 - ⇒ Enqueue any of these vertices whose In-Degree became zero



Sort this digraph!

Topological Sort Algorithm #2: Analysis

For input graph $G = (V, E)$, Run Time = ?

Break down into total time to:

Initialize In-Degree array:

$O(|E|)$

Initialize Queue with In-Degree 0 vertices:

$O(|V|)$

Dequeue and output vertex:

$|V|$ vertices, each takes only $O(1)$ to dequeue and output.

Total time = $O(|V|)$

Reduce In-Degree of all vertices adjacent to a vertex and


Enqueue any In-Degree 0 vertices:

$O(|E|)$

Total time = $O(|V| + |E|)$ **Linear running time!**

Paths

- ◆ Recall definition of a path in a tree – same for graphs
- ◆ A **path** is a list of vertices $\{v_1, v_2, \dots, v_n\}$ such that (v_i, v_{i+1}) is in E for all $0 \leq i < n$.



Example of a path:


$p = \{\text{Seattle, Salt Lake City, Chicago, Dallas, San Francisco, Seattle}\}$

Simple Paths and Cycles

- ◆ A **simple path** repeats no vertices (except the 1st can be the last):
 - ⇒ $p = \{\text{Seattle, Salt Lake City, San Francisco, Dallas}\}$
 - ⇒ $p = \{\text{Seattle, Salt Lake City, Dallas, San Francisco, Seattle}\}$
- ◆ A **cycle** is a path that starts and ends at the same node:
 - ⇒ $p = \{\text{Seattle, Salt Lake City, Dallas, San Francisco, Seattle}\}$
- ◆ A **simple cycle** is a cycle that repeats no vertices except that the first vertex is also the last
- ◆ A directed graph with no cycles is called a DAG (directed acyclic graph) E.g. All trees are DAGs
 - ⇒ A graph with cycles is often a drag...

Path Length and Cost

- ♦ **Path length**: the number of edges in the path
- ♦ **Path cost**: the sum of the costs of each edge
 - ⇒ Note: Path length = unweighted path cost (edge weight = 1)



Single Source, Shortest Path Problems


- ♦ Given a graph $G = (V, E)$ and a “source” vertex s in V , find the minimum cost paths from s to every vertex in V
- ♦ Many variations:
 - ⇒ unweighted vs. weighted
 - ⇒ cyclic vs. acyclic
 - ⇒ positive weights only vs. negative weights allowed
 - ⇒ multiple weight types to optimize
 - ⇒ Etc.
- ♦ We will look at only a couple of these...
 - ⇒ See text for the others

Why study shortest path problems?

- ◆ Plenty of applications
- ◆ **Traveling on a “starving student” budget:** What is the cheapest multi-stop airline schedule from Seattle to city X?
- ◆ **Optimizing routing of packets on the internet:**
 - ⇒ Vertices = routers, edges = network links with different delays
 - ⇒ *What is the routing path with smallest total delay?*
- ◆ **Hassle-free commuting:** Finding what highways and roads to take to minimize total delay due to traffic
- ◆ **Finding the fastest way to get to coffee vendors on campus from your classrooms**

Unweighted Shortest Paths Problem

Problem: Given a “source” vertex s in an unweighted graph $G = (V, E)$, find the shortest path from s to all vertices in G




Find the shortest path from C to: A B C D E F G H

Solution based on Breadth-First Search

- ♦ Basic Idea: Starting at node s , find vertices that can be reached using 0, 1, 2, 3, ..., $N-1$ edges (works even for cyclic graphs!)

On-board
example:




Find the shortest path from **C** to: **A B C D E F G H**

Breadth-First Search (BFS) Algorithm

- ♦ Uses a **queue** to store vertices that need to be expanded
- ♦ Pseudocode (source vertex is s):
 1. $\text{Dist}[s] = 0$
 2. $\text{Enqueue}(s)$
 3. **While** queue is not empty
 1. $X = \text{dequeue}$
 2. **For each** vertex Y adjacent to X and not previously visited
 - $\text{Dist}[Y] = \text{Dist}[X] + 1$ (Prev allows paths to be reconstructed)
 - $\text{Prev}[Y] = X$
 - $\text{Enqueue } Y$
- ♦ Running time (same as topological sort) = $O(|V| + |E|)$ (why?)

That was easy but what if edges have weights?

Does BFS still work for finding minimum cost paths?



Can you find a counterexample (a path) for this graph to show BFS won't work?

What if edges have weights?

- ♦ BFS does not work anymore – minimum cost path may have additional hops

Shortest path from

C to A:


BFS: C A

(cost = 9)

Minimum Cost

Path = C E D A

(cost = 8)



Dijkstra to the rescue...



E. W. Dijkstra
(1930-2002)

- ♦ Legendary figure in computer science
- ♦ Some rumors collected from previous classes...
- ♦ Rumor #1: Supported teaching introductory computer courses without computers (pencil and paper programming)
- ♦ Rumor #2: Supposedly wouldn't read his e-mail; so, his staff had to print out his e-mails and put them in his mailbox

An Aside: Dijkstra on GOTOs

“For a number of years I have been familiar with the observation that the **quality of programmers is a decreasing function of the density of go to statements** in the programs they produce.”

Opening sentence of: “*Go To Statement Considered Harmful*” by Edsger W. Dijkstra, Letter to the Editor, Communications of the ACM, Vol. 11, No. 3, March 1968, pp. 147-148.

Dijkstra's Algorithm for Weighted Shortest Path

- ♦ Classic algorithm for solving shortest path in weighted graphs (without negative weights)
- ♦ Example of a *greedy* algorithm
 - ⇒ Irrevocably makes decisions without considering future consequences
 - ⇒ Sound familiar? Not necessarily the best life strategy... but works in some cases (e.g. Huffman encoding)


Dijkstra's Algorithm for Weighted Shortest Path

- ♦ Basic Idea:
 - ⇒ Similar to BFS
 - ♦ Each vertex stores a cost for path from source
 - ♦ Vertex to be expanded is the one with least path cost seen so far
 - Greedy choice – always select current best vertex
 - Update costs of all neighbors of selected vertex
 - ⇒ But unlike BFS, a vertex already visited may be updated if a better path to it is found

Pseudocode for Dijkstra's Algorithm

1. Initialize the cost of each node to ∞
2. Initialize the cost of the source to 0
3. While there are unknown nodes left in the graph
 1. Select the unknown node N with the lowest cost (greedy choice)
 2. Mark N as known
 3. For each node X adjacent to N

If (N 's cost + cost of (N, X)) < X 's cost
 X 's cost = N 's cost + cost of (N, X)
 $Prev[X] = N$ //store preceding node



(Prev allows paths to be reconstructed)


Dijkstra's Algorithm (greed in action)

vertex	known	cost	Prev
A	No	∞	
B	No	∞	
C	Yes	0	
D	No	∞	
E	No	∞	

Initial

vertex	known	cost	Prev
A			
B			
C			
D			
E			

Final




Dijkstra's Algorithm (greed in action)

vertex	known	cost	Prev
A	No	∞	-
B	No	∞	-
C	Yes	0	-
D	No	∞	-
E	No	∞	-

Initial

vertex	known	cost	Prev
A	Yes	8	D
B	Yes	10	A
C	Yes	0	-
D	Yes	5	E
E	Yes	2	C

Final



Questions for Next Time:

Does Dijkstra's method always work?

How fast does it run?

Where else in life can I be greedy?

To Do:

Start Homework Assignment #4

(Don't wait until the last few days!!!)

Continue reading and enjoying chapter 9